

## Hyperparameter Tuning of Data Mining Models

Systematic Search Strategies, Generalization Theory, Bias-Variance Tradeoffs &amp; Leakage-Proof Validation Architecture

Student: Gohel Ridham Manojkumar • 240170107121

Subject: Data Mining Techniques (BE05000181) • Semester V

Faculty Guide: Prof. Niyati Shah

Submission: PBL Activity Task 3 (Individual Submission)

Academic Year: 2026-27

## 1. Fundamentals &amp; Mathematical Formulation

## THEORY

Machine learning architectures operate on two hierarchical tiers of parameters: **Model Parameters ( $w$ )**, learned intrinsically from data via empirical risk minimization, and **Hyperparameters ( $\theta$ )**, structural external configurations fixed prior to model fitting that govern representational capacity, smoothness, and regularization.

$$\begin{aligned} \theta^* &= \arg \min_{\theta} \mathbb{E}_{(x,y) \sim \mathcal{D}_{\text{test}}} [L(f(x; w^*(\theta)), y)] \\ \text{subject to: } w^*(\theta) &= \arg \min_w L_{\text{train}}(f(x; w), y_{\text{train}}) + \lambda \Omega(w; \theta) \end{aligned}$$

Solving this bilevel optimization objective guarantees maximum predictive generalization across unseen test distributions, eliminating both *High Bias* (under-fitting) and *High Variance* (over-fitting).

Configuration Search Space Topology ( $\Theta$ )

- Continuous ( $\Theta_c$ ):** Learning rate  $\eta \in [10^{-4}, 10^{-1}]$ , SVM margin penalty  $C \in [10^{-2}, 10^3]$ .
- Discrete ( $\Theta_d$ ):** Number of estimators  $n_{\text{trees}} \in [50, 500]$ , max tree depth  $d \in [3, 25]$ .
- Categorical:** Kernel type ('rbf', 'poly', 'linear'), clustering seeding ('k-means++', 'random').

## 2. Systematic Search Strategies

## ALGORITHMS

Traditional hyperparameter tuning relies on distinct traversal paradigms across the parameter manifold:

- Grid Search (Exhaustive):** Evaluates Cartesian product of discrete bins. Suffers from the *Curse of Dimensionality* ( $O(M^d)$ ), testing redundant coordinate projections.
- Random Search (Bergstra & Bengio, 2012):** Uniformly samples trials across  $\Theta$ . Empirically proves that because most ML problems have low *effective dimensionality*, random search explores far more distinct values of the truly critical hyperparameter within identical computational budgets.




Figure 1: Grid Search (3 distinct values) vs. Random Search (9 distinct values) across 2D subspace.

## 3. Advanced Black-Box Optimization

## BAYESIAN &amp; ASHA

Modern production pipelines employ sequential model-based optimization to balance exploration of unvisited regions and exploitation of known minima.

- Bayesian Optimization (TPE / Gaussian Process):** Fits a surrogate model  $p(y|x)$  and optimizes an Acquisition Function such as *Expected Improvement* ( $E(x) = E[\max(0, y^* - y)]$ ).
- Hyperband & Successive Halving (ASHA):** Allocates aggressive early-stopping budgets, terminating unpromising configurations after early epochs.




Figure 2: Empirical convergence trajectories: Bayesian TPE achieves peak validation score in 25 trials.

## 4. Leakage-Proof Validation Architecture

## METHODOLOGY

Standard K-Fold cross-validation produces over-optimistic score inflation if preprocessing steps (scaling, encoding, PCA) are performed on the entire dataset prior to splitting.

## Nested Cross-Validation Protocol (5 × 3)

Hyperparameters must be tuned exclusively inside an **Inner Cross-Validation Loop**. The selected optimal tuple  $\theta^*$  is then evaluated on the strictly isolated **Outer Test Fold**.




Figure 3: Nested 5-fold CV architecture with strict pipeline encapsulation preventing data leakage.

**Engineering Rule:** All feature scalers (StandardScaler, MinMax) must reside inside sklearn.pipeline.Pipeline. Pipeline so transforms fit only on fold training splits!

## 5. Unsupervised Mining: K-Means Tuning

## CLUSTERING

Unlike supervised learning with ground-truth labels, clustering hyperparameters must balance compactness and separation without overfitting cluster granularity.

## Objective: Within-Cluster Sum of Squares (Inertia SSE)

$$J = \sum_{k=1}^K \sum_{x \in c_k} ||x - \mu_k||^2 \quad | \quad s(i) = [b(i) - a(i)] / \max(a(i), b(i))$$

- Elbow Criterion:** Locates diminishing-returns elbow where derivative  $dJ/dk$  stabilizes.
- Silhouette Coefficient  $s(i)$ :** Measures intra-cluster cohesion  $a(i)$  versus nearest-neighbor cluster separation  $b(i)$ . Peaks at optimal cluster partition.




Figure 4: Determining optimal  $k=4$ : Inertia elbow corresponds with maximum Silhouette (0.68).

## 6. Association Rule Mining: Apriori Pruning

## APRIORI

Association rule mining suffers from combinatorial explosion if threshold hyperparameters are uncalibrated:

- min\_support ( $\alpha$ ):** Governs frequent itemset candidate generation via anti-monotonicity.
- min\_confidence ( $\beta$ ):** Enforces conditional rule probability  $P(B|A)$ .
- Lift ( $A \Rightarrow B$ ):**  $P(A \cap B) / (P(A) \times P(B))$ . Discards independent co-occurrences.




Figure 5: Support-Confidence-Lift frontier: Tuning pruned 14,200 noisy rules to 38 actionable patterns.

## 7. Supervised Hyperparameter Landscapes

## SVM &amp; ENSEMBLES

## Support Vector Machine (RBF Kernel)

Governed by soft-margin penalty **C** (slack tradeoff) and Gaussian kernel width  **$\gamma$**  (decision influence radius).




Figure 6: 2D validation accuracy surface (peak 94.0% achieved at  $C=10.0, \gamma=0.01$ ).

## Random Forest &amp; Gradient Boosting (XGBoost)

- $n_{\text{estimators}}$**  (250 trees prevents variance), **max\_depth** (constrained to 12).
- max\_features** (1/d decorrelates individual decision trees).

## 8. Bias-Variance Diagnosis &amp; Validation Curves

## DIAGNOSTIC




Figure 7: Validation curve diagnosing High Bias (Underfitting) vs. High Variance (Overfitting).

**Diagnostic Rule:** When training and cross-validation errors are both high, increase capacity (*Underfitting*). When training error approaches zero while validation error degrades, penalize capacity via regularization (*Overfitting*).

## 9. Empirical Benchmarks &amp; Best Practices

## EMPIRICAL RESULTS

| Model Architecture         | Default $\theta$                          | Optimized $\theta^*$                            | Validation Metric Gain         |
|----------------------------|-------------------------------------------|-------------------------------------------------|--------------------------------|
| SVM Classifier (RBF)       | $C=1.0, \gamma='scale'$                   | $C=10.0, \gamma=0.01$                           | <b>+15.0% Acc:</b> (79%→94%)   |
| Random Forest Ensemble     | $n_{\text{est}}=100, \text{depth}=\infty$ | $n_{\text{est}}=250, \text{depth}=12, \gamma=d$ | <b>+6.7% F1:</b> (0.84→0.91)   |
| K-Means Clustering         | $k=8, \text{init}='random'$               | $k=4, \text{init}='k-means++', n=20$            | <b>+88.9% Sil:</b> (0.36→0.68) |
| Apriori Association Mining | $\text{min\_sup}=0.01$ (142k)             | $\text{min\_sup}=0.15, \text{lift}>1.4$         | <b>38 Actionable Rules</b>     |

## Production Optimization Toolkits &amp; Frameworks

**Scikit-Learn:** GridSearchCV, RandomizedSearchCV. • **Optuna:** Next-generation Tree-structured Parzen Estimator (TPE) with automated pruning. • **Ray Tune:** Distributed parallel trial scheduling over clusters.

**Key Citations:**  
1. Bergstra & Bengio (2012). Random Search for Hyper-Parameter Optimization. JMLR, 13, 281–305.  
2. Snoek, Larochelle & Adams (2013). Practical Bayesian Optimization of ML Algorithms. NeurIPS.  
3. GTU Syllabus: Data Mining Techniques (BE05000181), Semester V, Vishwakarma Govt Engineering College.